The EXT2FS Library

by Theodore Ts'o

The EXT2FS Library

This manual documents the EXT2FS Library, version 1.42.12.

1 Introduction to the EXT2FS Library

The EXT2FS library is designed to allow user-level programs to manipulate an ext2 filesystem.

2 EXT2FS Library Functions

2.1 Filesystem-level functions

The following functions operate on a filesystem handle. Most EXT2FS Library functions require a filesystem handle as their first argument. There are two functions which create a filesystem handle, `ext2fs_open` and `ext2fs_initialize`.

The filesystem can also be closed using `ext2fs_close`, and any changes to the superblock and group descripts can be written out to disk using `ext2fs_flush`.

2.1.1 Opening an ext2 filesystem

Most libext2fs functions take a filesystem handle of type `ext2_filsys`. A filesystem handle is created either by opening an existing function using `ext2fs_open`, or by initializing a new filesystem using `ext2fs_initialize`.

errcode_t ext2fs_open (*const char* ***name**, *int* **flags**, *int*   [Function]
        superblock, *int* **block_size**, *io_manager* **manager**, *ext2_filsys* ***ret_fs**)
    Opens a filesystem named *name*, using the the io_manager *manager* to define the input/output routines needed to read and write the filesystem. In the case of the `unix_io` io_manager, *name* is interpreted as the Unix filename of the filesystem image. This is often a device file, such as '`/dev/hda1`'.

    The *superblock* parameter specifies the block number of the superblock which should be used when opening the filesystem. If *superblock* is zero, `ext2fs_open` will use the primary superblock located at offset 1024 bytes from the start of the filesystem image.

    The *block_size* parameter specifies the block size used by the filesystem. Normally this is determined automatically from the filesystem uperblock. If *block_size* is non-zero, it must match the block size found in the superblock, or the error `EXT2_ET_UNEXPECTED_BLOCK_SIZE` will be returned. The *block_size* parameter is also used to help fund the superblock when *superblock* is non-zero.

    The *flags* argument contains a bitmask of flags which control how the filesystem open should be handled.

    EXT2_FLAG_RW
            Open the filesystem for reading and writing. Without this flag, the filesystem is opened for reading only.

    EXT2_FLAG_FORCE
            Open the filesystem regardless of the feature sets listed in the superblock.

2.1.2 Closing and flushing out changes

errcode_t ext2fs_flush (*ext2_filsys* **fs**)   [Function]
    Write any changes to the high-level filesystem data structures in the *fs* filesystem. The following data structures will be written out:

- The filesystem superblock
- The filesystem group descriptors
- The filesystem bitmaps, if read in via `ext2fs_read_bitmaps`.

void ext2fs_free (*ext2 filsys* **fs**)   [Function]
>   Close the io_manager abstraction for *fs* and release all memory associated with the
>   filesystem handle.

errcode_t ext2fs_close (*ext2 filsys* **fs**)   [Function]
>   Flush out any changes to the high-level filesystem data structures using `ext2fs_flush`
>   if the filesystem is marked dirty; then close and free the filesystem using `ext2fs_free`.

2.1.3 Initializing a filesystem

An ext2 filesystem is initializing by the `mke2fs` program. The two functions described here,
`ext2fs_initialize` and `ext2fs_allocate_tables` do much of the initial work for setting
up a filesystem. However, they don't do the whole job. `mke2fs` calls `ext2fs_initialize` to
set up the filesystem superblock, and calls `ext2fs_allocate_tables` to allocate space for
the inode table, and the inode and block bitmaps. In addition, `mke2fs` must also initialize
the inode tables by clearing them with zeros, create the root and lost+found directories,
and reserve the reserved inodes.

errcode_t ext2fs_initialize (*const char *****name**, *int* **flags**, *struct*   [Function]
>   *ext2 super block ****param**, *io manager* **manager**, *ext2 filsys ****ret_fs**)
>   This function is used by the `mke2fs` program to initialize a filesystem. The `ext2fs_`
>   `initialize` function creates a filesystem handle which is returned in *ret fs* that has
>   been properly setup for a filesystem to be located in *name*, using the io_manager
>   *manager*. The prototype superblock in *param* is used to supply parameters such as
>   the number of blocks in the filesystem, the block size, etc.
>
>   The `ext2fs_initialize` function does not actually do any I/O; that will be done
>   when the application program calls `ext2fs_close` or `ext2fs_flush`. Also, this func-
>   tion only initializes the superblock and group descriptor structures. It does not create
>   the inode table or the root directory. This must be done by the calling application,
>   such as `mke2fs`.
>
>   The following values may be set in the *param* prototype superblock; if a value of 0 is
>   found in a field, `ext2fs_initialize` will use a default value. The calling application
>   should zero out the prototype entire superblock, and then fill in any appropriate
>   values.
>
>   `s_blocks_count`
>   The number of blocks in the filesystem. This parameter is mandatory
>   and must be set by the calling application.
>
>   `s_inodes_count`
>   The number of inodes in the filesystem. The default value is determined
>   by calculating the size of the filesystem, and creating one inode for every
>   4096 bytes.

s_r_blocks_count
> The number of blocks which should be reserved for the superuser. The default value is zero blocks.

s_log_block_size
> The blocksize of the filesystem. Valid values are 0 (1024 bytes), 1 (2048 bytes), or 2 (4096 bytes). The default blocksize is 1024 bytes.

s_log_frag_size
> The size of fragments. The ext2 filesystem does not support fragments (and may never support fragments). Currently this field must be the same as s_log_block_size.

s_first_data_block
> The first data block for the filesystem. For filesystem with a blocksize of 1024 bytes, this value must be at least 1, since the superblock is located in block number 1. For filesystems with larger blocksizes, the superblock is still located at an offset of 1024 bytes, so the superblock is located in block number 0. By default, this value is set to 1 for filesystems with a block size of 1024 bytes, or 0 for filesystems with larger blocksizes.

s_max_mnt_count
> This field defines the number of times that the filesystem can be mounted before it should be checked using e2fsck. When e2fsck is run without the '-f' option, e2fsck will skip the filesystem check if the number of times that the filesystem has been mounted is less than s_max_mnt_count and if the interval between the last time a filesystem check was performed and the current time is less than s_checkinterval (see below). The default value of s_max_mnt_count is 20.

s_checkinterval
> This field defines the minimal interval between filesystem checks. See the previous entry for a discussion of how this field is used by e2fsck. The default value of this field is 180 days (six months).

s_errors
> This field defines the behavior which should be used by the kernel of errors are detected in the filesystem. Possible values include:
>
> 'EXT2_ERRORS_CONTINUE'
>> Continue execution when errors are detected.
>
> 'EXT2_ERRORS_RO'
>> Remount the filesystem read-only.
>
> 'EXT2_ERRORS_PANIC'
>> Panic.
>
> The default behavior is 'EXT2_ERRORS_CONTINUE'.

errcode_t ext2fs_allocate_tables (ext2_filsys fs)   [Function]
> Allocate space for the inode table and the block and inode bitmaps. The inode tables and block and inode bitmaps aren't actually initialized; this function just allocates the space for them.

2.1.4 Filesystem flag functions

The filesystem handle has a number of flags which can be manipulated using the following function. Some of these flags affect how the libext2fs filesystem behaves; others are provided solely for the application's convenience.

void **ext2fs_mark_changed** (*ext2_filsys* **fs**)   [Function]
int **ext2fs_test_changed** (*ext2_filsys* **fs**)   [Function]

>   This flag indicates whether or not the filesystem has been changed. It is not used by the ext2fs library.

void **ext2fs_mark_super_dirty** (*ext2_filsys* **fs**)   [Function]

>   Mark the filesystem *fs* as being dirty; this will cause the superblock information to be flushed out when `ext2fs_close` is called. `ext2fs_mark_super_dirty` will also set the filesystem changed flag. The dirty flag is automatically cleared by `ext2fs_flush` when the superblock is written to disk.

void **ext2fs_mark_valid** (*ext2_filsys* **fs**)   [Function]
void **ext2fs_unmark_valid** (*ext2_filsys* **fs**)   [Function]
int **ext2fs_test_valid** (*ext2_filsys* **fs**)   [Function]

>   This flag indicates whether or not the filesystem is free of errors. It is not used by libext2fs, and is solely for the application's convenience.

void **ext2fs_mark_ib_dirty** (*ext2_filsys* **fs**)   [Function]
void **ext2fs_mark_bb_dirty** (*ext2_filsys* **fs**)   [Function]
int **ext2fs_test_ib_dirty** (*ext2_filsys* **fs**)   [Function]
int **ext2fs_test_bb_dirty** (*ext2_filsys* **fs**)   [Function]

>   These flags indicate whether or not the inode or block bitmaps have been modified. If the flag is set, it will cause the appropriate bitmap to be written when the filesystem is closed or flushed.

2.2 File I/O Functions

The following functions provide a convenient abstraction to read or write a file in an filesystem. The interface is similar in spirit to the Linux/POSIX file I/O system calls.

2.2.1 File handle manipulation

The file handle functions much like a file descriptor in the Linux/POSIX file I/O system calls. Unlike the Linux/POSIX system calls, files are opened via inode numbers instead of via pathnames. To resolve a pathname to an inode number, use the function `ext2fs_namei` or to create a new file, use `ext2fs_new_inode` and `ext2fs_link`.

errcode_t **ext2fs_file_open2** (*ext2_filsys* **fs**, *ext2_ino_t* **ino**, *struct*   [Function]
        ext2_inode ***inode**, *int* **flags**, *ext2_file_t* ***ret**)
errcode_t **ext2fs_file_open** (*ext2_filsys* **fs**, *ext2_ino_t* **ino**, *int*   [Function]
        flags, *ext2_file_t* ***ret**)

>   Opens a file identified by inode number *ino* in filesystem *fs* and returns a file handle in *ret*. If an inode structure is provided in *inode*, then it is used instead of reading the inode from the filesystem.

The *flags* argument contains a bitmask of flags which control how the file should be opened.

EXT2_FILE_WRITE
> Open the file for reading and writing. Without this flag, the file is opened for writing only.

EXT2_FILE_CREATE
> Create the file if it is not already present.

ext2_filsys **ext2fs_file_get_fs** (*ext2_file_t* `file`)   [Function]
> Return the filesystem handle where the open file *file* was opened.

errcode_t **ext2fs_file_close** (*ext2_file_t* `file`)   [Function]
> Close the file handle *file*.

errcode_t **ext2fs_file_flush** (*ext2_file_t* `file`)   [Function]
> Force any data written via `ext2fs_file_write` to disk.

2.2.2 Reading and writing data

errcode_t **ext2fs_file_read** (*ext2_file_t* `file`, *void *`buf`, *unsigned*   [Function]
> *int* `wanted`, *unsigned int *`got`)
> Read *wanted* bytes of data from *file* store it in the buffer *buf*. The number of bytes that was actually read is returned via *got*.

errcode_t **ext2fs_file_write** (*ext2_file_t* `file`, *const void *`buf`,   [Function]
> *unsigned int* `nbytes`, *unsigned int *`written`)
> Write *wanted* bytes of data from the buffer *buf* to the current file position of *file*. The number of bytes that was actually written is returned via *got*.

2.2.3 Changing the file offset

errcode_t **ext2fs_file_llseek** (*ext2_file_t* `file`, *__u64* `offset`, *int*   [Function]
> *whence*, *__u64 *`ret_pos`)
errcode_t **ext2fs_file_lseek** (*ext2_file_t* `file`, *ext2_off_t* `offset`,   [Function]
> *int* `whence`, *ext2_off_t *`ret_pos`)
> Change the current file position of *file* according to the directive *whence* as follows:

EXT2_SEEK_SET
> The file position is set to *offset* bytes from the beginning of the file.

EXT2_SEEK_CUR
> The file position set to its current location plus *offset* bytes.

EXT2_SEEK_END
> The file position is set to the size of the file plus *offset* bytes.

The current offset is returned via *ret_pos*.

2.2.4 Getting the file size

`errcode_t ext2fs_file_get_lsize` (*ext2_file_t* `file`, *__u64*   [Function]
          *`*ret_size`)
>   Return the size of the file *file* in *ret_size*.

`ext2_off_t ext2fs_file_get_size` (*ext2_file_t* `file`)   [Function]
>   Return the size of the file *file*.

2.3 Inode Functions

2.3.1 Reading and writing inodes

`errcode_t ext2fs_read_inode` (*ext2_filsys* `fs`, *ext2_ino_t* `ino`, *struct*   [Function]
          ext2_inode `*inode`)
>   Read the inode number *ino* into *inode*.

`errcode_t ext2fs_write_inode` (*ext2_filsys* `fs`, *ext2_ino_t* `ino`, *struct*   [Function]
          ext2_inode `*inode`)
>   Write *inode* to inode *ino*.

2.3.2 Iterating over inodes in a filesystem

The inode_scan abstraction is useful for iterating over all the inodes in a filesystem.

`errcode_t ext2fs_open_inode_scan` (*ext2_filsys* `fs`, *int*   [Function]
          `buffer_blocks`, *ext2_inode_scan* `*scan`)
>   Initialize the iteration variable *scan*. This variable is used by `ext2fs_get_next_`
>   `inode`. The *buffer_blocks* parameter controls how many blocks of the inode table are
>   read in at a time. A large number of blocks requires more memory, but reduces the
>   overhead in seeking and reading from the disk. If *buffer_blocks* is zero, a suitable
>   default value will be used.

`void ext2fs_close_inode_scan` (*ext2_inode_scan* `scan`)   [Function]
>   Release the memory associated with *scan* and invalidate it.

`errcode_t ext2fs_get_next_inode` (*ext2_inode_scan* `scan`, *ext2_ino_t*   [Function]
          `*ino`, *struct ext2_inode* `*inode`)
>   This function returns the next inode from the filesystem; the inode number of the
>   inode is stored in *ino*, and the inode is stored in *inode*.
>
>   If the inode is located in a block that has been marked as bad, `ext2fs_get_next_`
>   `inode` will return the error `EXT2_ET_BAD_BLOCK_IN_INODE_TABLE`.

`errcode_t ext2fs_inode_scan_goto_blockgroup` (*ext2_inode_scan*   [Function]
          `scan`, *int* `group`)
>   Start the inode scan at a particular ext2 blockgroup, *group*. This function may be
>   safely called at any time while *scan* is valid.

void **ext2fs_set_inode_callback** (*ext2_inode_scan* `scan`, *errcode_t*   [Function]
     (**done_group*)(*ext2_filsys* `fs`, *ext2_inode_scan* `scan`, *dgrp_t* `group`, *void **
     `private`), *void **`done_group_data`)
> Register a callback function which will be called by `ext2_get_next_inode` when all
> of the inodes in a block group have been processed.

int **ext2fs_inode_scan_flags** (*ext2_inode_scan* `scan`, *int*   [Function]
     `set_flags`, *int* `clear_flags`)
> Set the scan_flags *set_flags* and clear the scan_flags *clear_flags*. The following flags
> can be set using this interface:

> `EXT2_SF_SKIP_MISSING_ITABLE`
>> When a block group is missing an inode table, skip it. If this
>> flag is not set `ext2fs_get_next_inode` will return the error
>> EXT2_ET_MISSING_INODE_TABLE.

2.3.3 Iterating over blocks in an inode

errcode_t **ext2fs_block_iterate** (*ext2_filsys* `fs`, *ext2_ino_t* `ino`, *int*   [Function]
     `flags`, *char *block_buf*, *int* (**func*)(*ext2_filsys* `fs`, *blk_t **`blocknr`, *int*
     `blockcnt`, *void **`private`), *void **`private`)
> Iterate over all of the blocks in inode number *ino* in filesystem *fs*, by calling the
> function *func* for each block in the inode. The *block_buf* parameter should either be
> NULL, or if the `ext2fs_block_iterate` function is called repeatedly, the overhead
> of allocating and freeing scratch memory can be avoided by passing a pointer to a
> scratch buffer which must be at least as big as three times the filesystem's blocksize.

> The *flags* parameter controls how the iterator will function:

> `BLOCK_FLAG_HOLE`
>> This flag indiciates that the interator function should be called on blocks
>> where the block number is zero (also known as "holes".) It is also known
>> as BLOCK_FLAG_APPEND, since it is also used by functions such as
>> ext2fs_expand_dir() to add a new block to an inode.

> `BLOCK_FLAG_DEPTH_TRAVERSE`
>> This flag indicates that the iterator function for the indirect, doubly in-
>> direct, etc. blocks should be called after all of the blocks containined in
>> the indirect blocks are processed. This is useful if you are going to be
>> deallocating blocks from an inode.

> `BLOCK_FLAG_DATA_ONLY`
>> This flag indicates that the iterator function should be called for data
>> blocks only.

> The callback function *func* is called with a number of parameters; the *fs* and *pri-*
> *vate* parameters are self-explanatory, and their values are taken from the parameters
> to `ext2fs_block_iterate`. (The *private* data structure is generally used by callers
> to `ext2fs_block_iterate` so that some private data structure can be passed to

the callback function. The *blockcnt* parameter, if non-negative, indicates the logi-
cal block number of a data block in the inode. If *blockcnt* is less than zero, then
func was called on a metadata block, and *blockcnt* will be one of the following val-
ues: BLOCK_COUNT_IND, BLOCK_COUNT_DIND, BLOCK_COUNT_TIND, or
BLOCK_COUNT_TRANSLATOR. The *blocknr* is a pointer to the inode or indirect
block entry listing physical block number. The callback function may modify the
physical block number, if it returns the *BLOCK_CHANGED* flag.

The callback function *func* returns a result code which is composed of the logical OR
of the following flags:

'`BLOCK_CHANGED`'
> This flag indicates that callback function has modified the physical block
> number pointed to by *blocknr*.

'`BLOCK_ABORT`'
> This flag requests that `ext2fs_block_iterate` to stop immediately and
> return to the caller.

errcode_t **ext2fs_block_iterate2** (*ext2_filsys* **fs**, *ext2_ino_t* **ino**, *int*   [Function]
        flags**, *char* ***block**_buf, *int* (**func**)(*ext2_filsys* **fs**, *blk_t* ***blocknr**,
        e2_blkcnt_t **blockcnt**, *blk_t* **ref_blk**, *int* **ref_offset**, *void* ***private**), *void*
        ***private**)
        This function is much like `ext2fs_block_iterate`, except that the *blockcnt* type is
        a 64-bit signed quantity, to support larger files, and the addition of the *ref_blk* and
        ref_offset arguments passed to the callback function, which identify the location of
        the physical block pointed to by pointer *blocknr*. If *ref_blk* is zero, then *ref_offset*
        contains the offset into the `i_blocks` array. If *ref_blk* is non-zero, then the physical
        block location is contained inside an indirect block group, and *ref_offset* contains the
        offset into the indirect block.

2.3.4 Convenience functions for Inodes

errcode_t **ext2fs_get_blocks** (*ext2_filsys* **fs**, *ext2_ino_t* **ino**, *blk_t*   [Function]
        ***blocks**)
        Returns an array of blocks corresponding to the direct, indirect, doubly indirect, and
        triply indirect blocks as stored in the inode structure.

errcode_t **ext2fs_check_directory** (*ext2_filsys* **fs**, *ext2_ino_t* **ino**)   [Function]
        Returns 0 if *ino* is a directory, and `ENOTDIR` if it is not.

int **ext2fs_inode_has_valid_blocks** (*struct ext2_inode* ***inode**)   [Function]
        Returns 1 if the inode's block entries actually valid block entries, and 0 if not. Inodes
        which represent devices and fast symbolic links do not contain valid block entries.

2.4 Directory functions

2.4.1 Directory block functions

errcode_t ext2fs_read_dir_block (*ext2_filsys* **fs**, *blk_t* **block**, *void*   [Function]
        *buf)
>   This function reads a directory block, performing any necessary byte swapping if
>   necessary.

errcode_t ext2fs_write_dir_block (*ext2_filsys* **fs**, *blk_t* **block**,   [Function]
        void *buf)
>   This function writes a directory block, performing any necessary byte swapping if
>   necessary.

errcode_t ext2fs_new_dir_block (*ext2_filsys* **fs**, *ext2_ino_t*   [Function]
        dir_ino, *ext2_ino_t* **parent_ino**, *char **block*)
>   This function creates a new directory block in *block*. If *dir_ino* is non-zero, then
>   *dir_info* and *parent_ino* is used to initialize directory entries for '.' and '..', respec-
>   tively.

2.4.2 Iterating over a directory

errcode_t ext2fs_dir_iterate (*ext2_filsys* **fs**, *ext2_ino_t* **dir**, *int*   [Function]
        flags, *char *block_buf*, *int* (**func*)(*struct ext2_dir_entry *dirent*, *int*
        offset, *int* **blocksize**, *char *buf*, *void *private*), *void *private*)
>   This function iterates over all of the directory entries in the directory *dir*, calling
>   the callback function *func* for each directory entry in the directory. The *block_buf*
>   parameter should either be NULL, or if the `ext2fs_dir_iterate` function is called
>   repeatedly, the overhead of allocating and freeing scratch memory can be avoided by
>   passing a pointer to a scratch buffer which must be at least as big as the filesystem's
>   blocksize.
>
>   The *flags* parameter controls how the iterator will function:
>
>   'DIRENT_FLAG_INCLUDE_EMPTY'
>   This flag indicates that the callback function should be called even for
>   deleted or empty directory entries.

2.4.3 Creating and expanding directories

errcode_t ext2fs_mkdir (*ext2_filsys* **fs**, *ext2_ino_t* **parent**, *ext2_ino_t*   [Function]
        inum, *const char *name*)
>   This function creates a new directory. If *inum* is zero, then a new inode will be
>   allocated; otherwise, the directory will be created in the inode specified by *inum*.
>   If *name* specifies the name of the new directory; if it is non-NULL, then the new
>   directory will be linked into the parent directory *parent*.

errcode_t ext2fs_expand_dir (*ext2_filsys* **fs**, *ext2_ino_t* **dir**)   [Function]
>   This function adds a new empty directory block and appends it to the directory *dir*.
>   This allows functions such as `ext2fs_link` to add new directory entries to a directory
>   which is full.

2.4.4 Creating and removing directory entries

errcode_t ext2fs_link (*ext2_filsys* **fs**, *ext2_ino_t* **dir**, *const char*   [Function]
          ***name**, *ext2_ino_t* **ino**, *int flags*)
> This function adds a new directory entry to the directory *dir*, with *name* and *ino*
> specifying the name and inode number in the directory entry, respectively.
>
> The low 3 bits of the flags field is used to specify the file type of inode: (No other
> flags are currently defined.)
>
> 'EXT2_FT_UNKNOWN'
>   The file type is unknown.
>
> 'EXT2_FT_REG_FILE'
>   The file type is a normal file.
>
> 'EXT2_FT_DIR'
>   The file type is a directory.
>
> 'EXT2_FT_CHRDEV'
>   The file type is a character device.
>
> 'EXT2_FT_BLKDEV'
>   The file type is a block device.
>
> 'EXT2_FT_FIFO'
>   The file type is a named pipe.
>
> 'EXT2_FT_SOCK'
>   The file type is a unix domain socket.
>
> 'EXT2_FT_SYMLINK'
>   The file type is a symbolic link.

errcode_t ext2fs_unlink (*ext2_filsys* **fs**, *ext2_ino_t* **dir**, *const char*   [Function]
          ***name**, *ext2_ino_t* **ino**, *int flags*)
> This function removes a directory entry from *dir*. The directory entry to be removed
> is the first one which is matched by *name* and *ino*. If *name* is non-NULL, the directory
> entry's name must match *name*. If *ino* is non-zero, the directory entry's inode number
> must match *ino*. No flags are currently defined for ext2fs_unlink; callers should pass
> in zero to this parameter.

2.4.5 Looking up filenames

errcode_t ext2fs_lookup (*ext2_filsys* **fs**, *ext2_ino_t* **dir**, *const char*   [Function]
          ***name**, *int* **namelen**, *char* ***buf**, *ext2_ino_t* ***inode**)

errcode_t ext2fs_namei (*ext2_filsys* **fs**, *ext2_ino_t* **root**, *ext2_ino_t*   [Function]
          cwd, *const char* ***name**, *ext2_ino_t* ***inode**)

errcode_t ext2fs_namei_follow (*ext2_filsys* **fs**, *ext2_ino_t* **root**,   [Function]
          ext2_ino_t **cwd**, *const char* ***name**, *ext2_ino_t* ***inode**)

errcode_t ext2fs_follow_link (*ext2_filsys* **fs**, *ext2_ino_t* **root**,   [Function]
          ext2_ino_t **cwd**, *ext2_ino_t* **inode**, *ext2_ino_t* ***res**_inode)

2.4.6 Translating inode numbers to filenames

errcode_t ext2fs_get_pathname (*ext2_filsys* `fs`, *ext2_ino_t* `dir`,   [Function]
     ext2_ino_t `ino`, *char ***`name`)

2.5 Bitmap Functions

2.5.1 Reading and Writing Bitmaps

errcode_t ext2fs_write_inode_bitmap (*ext2_filsys* `fs`)   [Function]

errcode_t ext2fs_write_block_bitmap (*ext2_filsys* `fs`)   [Function]

errcode_t ext2fs_read_inode_bitmap (*ext2_filsys* `fs`)   [Function]

errcode_t ext2fs_read_block_bitmap (*ext2_filsys* `fs`)   [Function]

errcode_t ext2fs_read_bitmaps (*ext2_filsys* `fs`)   [Function]

errcode_t ext2fs_write_bitmaps (*ext2_filsys* `fs`)   [Function]

2.5.2 Allocating Bitmaps

errcode_t ext2fs_allocate_generic_bitmap (*__u32* `start`, *__u32*   [Function]
    `end`, *_u32* `real_end`, *const char **`descr`, *ext2fs_generic_bitmap **`ret`)

errcode_t ext2fs_allocate_block_bitmap (*ext2_filsys* `fs`, *const*   [Function]
    char *`descr`, *ext2fs_block_bitmap **`ret`)

errcode_t ext2fs_allocate_inode_bitmap (*ext2_filsys* `fs`, *const*   [Function]
    char *`descr`, *ext2fs_inode_bitmap **`ret`)

2.5.3 Freeing bitmaps

void ext2fs_free_generic_bitmap (*ext2fs_inode_bitmap* `bitmap`)   [Function]

void ext2fs_free_block_bitmap (*ext2fs_block_bitmap* `bitmap`)   [Function]

void ext2fs_free_inode_bitmap (*ext2fs_inode_bitmap* `bitmap`)   [Function]

2.5.4 Bitmap Operations

void ext2fs_mark_block_bitmap (*ext2fs_block_bitmap* `bitmap`, *blk_t*   [Function]
    `block`)

void ext2fs_unmark_block_bitmap (*ext2fs_block_bitmap* `bitmap`,   [Function]
    blk_t `block`)

int ext2fs_test_block_bitmap (*ext2fs_block_bitmap* `bitmap`, *blk_t*   [Function]
    `block`)

    These functions set, clear, and test bits in a block bitmap *bitmap*.

void ext2fs_mark_inode_bitmap (*ext2fs_inode_bitmap* `bitmap`,   [Function]
    ext2_ino_t `inode`)

void ext2fs_unmark_inode_bitmap (*ext2fs_inode_bitmap* `bitmap`,   [Function]
    ext2_ino_t `inode`)

int ext2fs_test_inode_bitmap (*ext2fs_inode_bitmap* `bitmap`,   [Function]
    ext2_ino_t `inode`)

    These functions set, clear, and test bits in an inode bitmap *bitmap*.

void ext2fs_fast_mark_block_bitmap (*ext2fs_block_bitmap* `bitmap`,   [Function]
        blk_t `block`)

void ext2fs_fast_unmark_block_bitmap (*ext2fs_block_bitmap*   [Function]
        `bitmap`, *blk_t* `block`)

int ext2fs_fast_test_block_bitmap (*ext2fs_block_bitmap* `bitmap`,   [Function]
        blk_t `block`)

void ext2fs_fast_mark_inode_bitmap (*ext2fs_inode_bitmap* `bitmap`,   [Function]
        ext2_ino_t `inode`)

void ext2fs_fast_unmark_inode_bitmap (*ext2fs_inode_bitmap*   [Function]
        `bitmap`, *ext2_ino_t* `inode`)

int ext2fs_fast_test_inode_bitmap (*ext2fs_inode_bitmap* `bitmap`,   [Function]
        ext2_ino_t `inode`)

> These "fast" functions are like their normal counterparts; however, they are imple-
> mented as inline functions and do not perform bounds checks on the inode number
> or block number; they are assumed to be correct. They should only be used in speed-
> critical applications, where the inode or block number has already been validated by
> other means.

blk_t ext2fs_get_block_bitmap_start (*ext2fs_block_bitmap*   [Function]
        `bitmap`)

ext2_ino_t ext2fs_get_inode_bitmap_start (*ext2fs_inode_bitmap*   [Function]
        `bitmap`)

> Return the first inode or block which is stored in the bitmap.

blk_t ext2fs_get_block_bitmap_end (*ext2fs_block_bitmap* `bitmap`)   [Function]

ext2_ino_t ext2fs_get_inode_bitmap_end (*ext2fs_inode_bitmap*   [Function]
        `bitmap`)

> Return the last inode or block which is stored in the bitmap.

2.5.5 Comparing bitmaps

errcode_t ext2fs_compare_block_bitmap (*ext2fs_block_bitmap* `bm1`,   [Function]
        ext2fs_block_bitmap `bm2`)

errcode_t ext2fs_compare_inode_bitmap (*ext2fs_inode_bitmap* `bm1`,   [Function]
        ext2fs_inode_bitmap `bm2`)

2.5.6 Modifying Bitmaps

errcode_t ext2fs_fudge_inode_bitmap_end (*ext2fs_inode_bitmap*   [Function]
        `bitmap`, *ext2_ino_t* `end`, *ext2_ino_t* `*oend`)

errcode_t ext2fs_fudge_block_bitmap_end (*ext2fs_block_bitmap*   [Function]
        `bitmap`, *blk_t* `end`, *blk_t* `*oend`)

2.5.7 Resizing Bitmaps

errcode_t ext2fs_resize_generic_bitmap (*__u32* `new_end`, *__u32*   [Function]
        `new_real_end`, *ext2fs_generic_bitmap* `bmap`)

errcode_t ext2fs_resize_inode_bitmap (*__u32* `new_end`, *__u32*   [Function]
        `new_real_end`, *ext2fs_inode_bitmap* `bmap`)

errcode_t ext2fs_resize_block_bitmap (__u32 new_end, __u32   [Function]
        new_real_end, ext2fs_block_bitmap bmap)

2.5.8 Clearing Bitmaps

void ext2fs_clear_inode_bitmap (ext2fs_inode_bitmap bitmap)   [Function]
        This function sets all of the bits in the inode bitmap bitmap to be zero.

void ext2fs_clear_block_bitmap (ext2fs_block_bitmap bitmap)   [Function]
        This function sets all of the bits in the block bitmap bitmap to be zero.

2.6 EXT2 data abstractions

The ext2 library has a number of abstractions which are useful for ext2 utility programs.

2.6.1 Badblocks list management

errcode_t ext2fs_badblocks_list_create (ext2_badblocks_list   [Function]
        *ret, int size)

void ext2fs_badblocks_list_free (ext2_badblocks_list bb)   [Function]

errcode_t ext2fs_badblocks_list_add (ext2_badblocks_list bb,   [Function]
        blk_t blk)

int ext2fs_badblocks_list_test (ext2_badblocks_list bb, blk_t blk) [Function]

errcode_t ext2fs_badblocks_list_iterate_begin   [Function]
        (ext2_badblocks_list bb, ext2_badblocks_iterate *ret)

int ext2fs_badblocks_list_iterate (ext2_badblocks_iterate iter,   [Function]
        blk_t *blk)

void ext2fs_badblocks_list_iterate_end (ext2_badblocks_iterate   [Function]
        iter)

errcode_t ext2fs_update_bb_inode (ext2_filsys fs,   [Function]
        ext2_badblocks_list bb_list)

errcode_t ext2fs_read_bb_inode (ext2_filsys fs, ext2_badblocks_list [Function]
        *bb_list)

errcode_t ext2fs_read_bb_FILE (ext2_filsys fs, FILE *f,   [Function]
        ext2_badblocks_list *bb_list, void (*invalid)(ext2_filsys fs, blk_t blk))

2.6.2 Directory-block list management

The dblist abstraction stores a list of blocks belonging to directories. This list can be useful
when a program needs to interate over all directory entries in a filesystem; e2fsck does this
in pass 2 of its operations, and debugfs needs to do this when it is trying to turn an inode
number into a pathname.

errcode_t ext2fs_init_dblist (ext2_filsys fs, ext2_dblist   [Function]
        *ret_dblist)
        Creates a dblist data structure and return it in ret_dblist.

void ext2fs_free_dblist (*ext2_dblist* `dblist`)   [Function]
> Free a dblist data structure.

errcode_t ext2fs_add_dir_block (*ext2_dblist* `dblist`, *ext2_ino_t*   [Function]
> `ino`, *blk_t* `blk`, *int* `blockcnt`)
> Add an entry to the dblist data structure. This call records the fact that block number
> *blockcnt* of directory inode *ino* is stored in block *blk*.

errcode_t ext2fs_set_dir_block (*ext2_dblist* `dblist`, *ext2_ino_t*   [Function]
> `ino`, *blk_t* `blk`, *int* `blockcnt`)
> Change an entry in the dblist data structure; this changes the location of block number
> *blockcnt* of directory indoe *ino* to be block *blk*.

errcode_t ext2fs_dblist_iterate (*ext2_dblist* `dblist`, *int*   [Function]
> (*func)(*ext2_filsys* `fs`, *struct ext2_db_entry* *`db_info`, *void* *`private`), *void*
> *`private`)
> This iterator calls *func* for every entry in the dblist data structure.

errcode_t ext2fs_dblist_dir_iterate (*ext2_dblist* `dblist`, *int*   [Function]
> *flags*, *char* *`block_buf`, *int* (*func)(*ext2_ino_t* `dir`, *int* `entry`, *struct*
> *ext2_dir_entry* *`dirent`, *int* `offset`, *int* `blocksize`, *char* *`buf`, *void*
> *`private`), *void* *`private`)
> This iterator takes reads in the directory block indicated in each dblist entry, and
> calls *func* for each directory entry in each directory block. If *dblist* contains all the
> directory blocks in a filesystem, this function provides a convenient way to iterate
> over all directory entries for that filesystem.

2.6.3 Inode count functions

The icount abstraction is a specialized data type used by `e2fsck` to store how many times a
particular inode is referenced by the filesystem. This is used twice; once to store the actual
number of times that the inode is reference; and once to store the claimed number of times
the inode is referenced according to the inode structure.

This abstraction is designed to be extremely efficient for storing this sort of information,
by taking advantage of the following properties of inode counts, namely (1) inode counts
are very often zero (because the inode is currrently not in use), and (2) many files have a
inode count of 1 (because they are a file which has no additional hard links).

errcode_t ext2fs_create_icount2 (*ext2_filsys* `fs`, *int* `flags`, *int*   [Function]
> `size`, *ext2_icount_t* `hint`, *ext2_icount_t* *`ret`)
> Creates an icount stucture for a filesystem *fs*, with initial space for *size* inodes
> whose count is greater than 1. The *flags* parameter is either 0 or `EXT2_ICOUNT_`
> `OPT_INCREMENT`, which indicates that icount structure should be able to increment
> inode counts quickly. The icount structure is returned in *ret*. The returned icount
> structure initially has a count of zero for all inodes.
>
> The *hint* parameter allows the caller to optionally pass in another icount structure
> which is used to initialize the array of inodes whose count is greater than 1. It is used

purely as a speed optimization so that the icount structure can determine in advance which inodes are likely to contain a count grater than 1.

void ext2fs_free_icount (*ext2_icount_t* `icount`)   [Function]
    Frees an icount structure.

errcode_t ext2fs_icount_fetch (*ext2_icount_t* `icount`, *ext2_ino_t*   [Function]
      ino, *__u16 *ret*)
    Returns in *ret* fetches the count for a particular inode *ino*.

errcode_t ext2fs_icount_increment (*ext2_icount_t* `icount`,   [Function]
      ext2_ino_t `ino`, *__u16 *ret*)
    Increments the ref count for inode *ino*.

errcode_t ext2fs_icount_decrement (*ext2_icount_t* `icount`,   [Function]
      ext2_ino_t `ino`, *__u16 *ret*)
    Decrements the ref count for inode *ino*.

errcode_t ext2fs_icount_store (*ext2_icount_t* `icount`, *ext2_ino_t*   [Function]
      ino, *__u16 count*)
    Sets the reference count for inode *ino* to be *count*.

ext2_ino_t ext2fs_get_icount_size (*ext2_icount_t* `icount`)   [Function]
    Returns the current number of inodes in *icount* which has a count greater than 1.

errcode_t ext2fs_icount_validate (*ext2_icount_t* `icount`, *FILE*   [Function]
      *f)
    Validates the internal rep invariant of *icount*; if there are any problems, print out debugging information to *f*. This function is intended for debugging and testing use only.

2.7 Byte-swapping functions

void ext2fs_swap_super (*struct ext2_super_block * super*)   [Function]

void ext2fs_swap_group_desc (*struct ext2_group_desc *gdp*)   [Function]

void ext2fs_swap_inode (*ext2_filsys* `fs`, *struct ext2_inode *to*, *struct*   [Function]
      ext2_inode *from, *int* `hostorder`)

int ext2fs_native_flag (*void*)   [Function]

2.8 Other functions

/* alloc.c */

errcode_t ext2fs_new_inode (*ext2_filsys* `fs`, *ext2_ino_t* `dir`, *int* `mode`,   [Function]
        ext2fs_inode_bitmap `map`, *ext2_ino_t* `*ret`)

errcode_t ext2fs_new_block (*ext2_filsys* `fs`, *blk_t* `goal`,   [Function]
        ext2fs_block_bitmap `map`, *blk_t* `*ret`)

errcode_t ext2fs_get_free_blocks (*ext2_filsys* `fs`, *blk_t* `start`,   [Function]
        blk_t `finish`, *int* `num`, *ext2fs_block_bitmap* `map`, *blk_t* `*ret`)
/* check_desc.c */

errcode_t ext2fs_check_desc (*ext2_filsys* `fs`)   [Function]

errcode_t ext2fs_get_num_dirs (*ext2_filsys* `fs`, *ext2_ino_t*   [Function]
        `*ret_num_dirs`)
/* getsize.c */

errcode_t ext2fs_get_device_size (*const char* `*file`, *int*   [Function]
        `blocksize`, *blk_t* `*retblocks`)
/* ismounted.c */

errcode_t ext2fs_check_if_mounted (*const char* `*file`, *int*   [Function]
        `*mount_flags`)
/* version.c */

int ext2fs_get_library_version (*const char* `**ver_string`, *const*   [Function]
        char `**date_string`)
    This function returns the current version of the ext2 library. The return value contains
    an integer version code, which consists of the major version number of the library
    multiplied by 100, plus the minor version number of the library. Hence, if the library
    version is 1.08, the returned value will be 108.

    If *ver_string* and/or *date_string* are non-NULL, they will be set to point at a constant
    string containing the library version and/or release date, respectively.

int ext2fs_parse_version_string (*const char* `*ver_string`)   [Function]
    This function takes a version string which may included in an application and returns
    a version code using the same algorithm used by `ext2fs_get_library_version`. It
    can be used by programs included in the `e2fsprogs` distribution to assure that they
    are using an up-to-date ext2 shared library.

/* inline functions */

int ext2fs_group_of_blk (*ext2_filsys* `fs`, *blk_t* `blk`)   [Function]
    This function returns the block group which contains the block *blk*.

int ext2fs_group_of_ino (*ext2_filsys* `fs`, *ext2_ino_t* `ino`)   [Function]
    This function returns the block group which contains the inode *ino*.

Concept Index

(Index is nonexistent)

Function and Type Index

Table of Contents